

# Machine Learning Classifier Ensembles

- Ensemble types: homogeneous vs heterogeneous
- Weighting schemes
- Ensemble schemes
  - Bagging
  - Boosting
  - Stacking
  - Classifiers:
    - Random Forest/Rotation Forest
    - Adaboost/XGBoost

arXiv.org > cs > arXiv:1809.06705

Computer Science > Machine Learning

Is rotation forest the best classifier for problems with continuous features?

A. Bagnall, M. Flynn, J. Large, J. Line, A. Bostrom, G. Cawley  
(Submitted on 18 Sep 2018 (v1), last revised 12 Sep 2019 (this version, v2))

Search... All fields Help | Advanced Search

Download:

- PDF
- Other formats (license)

Current browse context: cs.LG  
< prev | next >

Open Access | Published: 17 June 2019

A probabilistic classifier ensemble weighting scheme based on cross-validated accuracy estimates

James Large, Jason Lines & Anthony Bagnall

*Data Mining and Knowledge Discovery* 33, 1674–1709(2019) | Cite this article

2140 Accesses | 4 Citations | 1 Altmetric | Metrics

Download PDF

Sections

Figures

References

Abstract

Introduction

Background

The cross-validation accuracy weighted pr.

# Ensembles

Different classifiers make different kind of mistakes

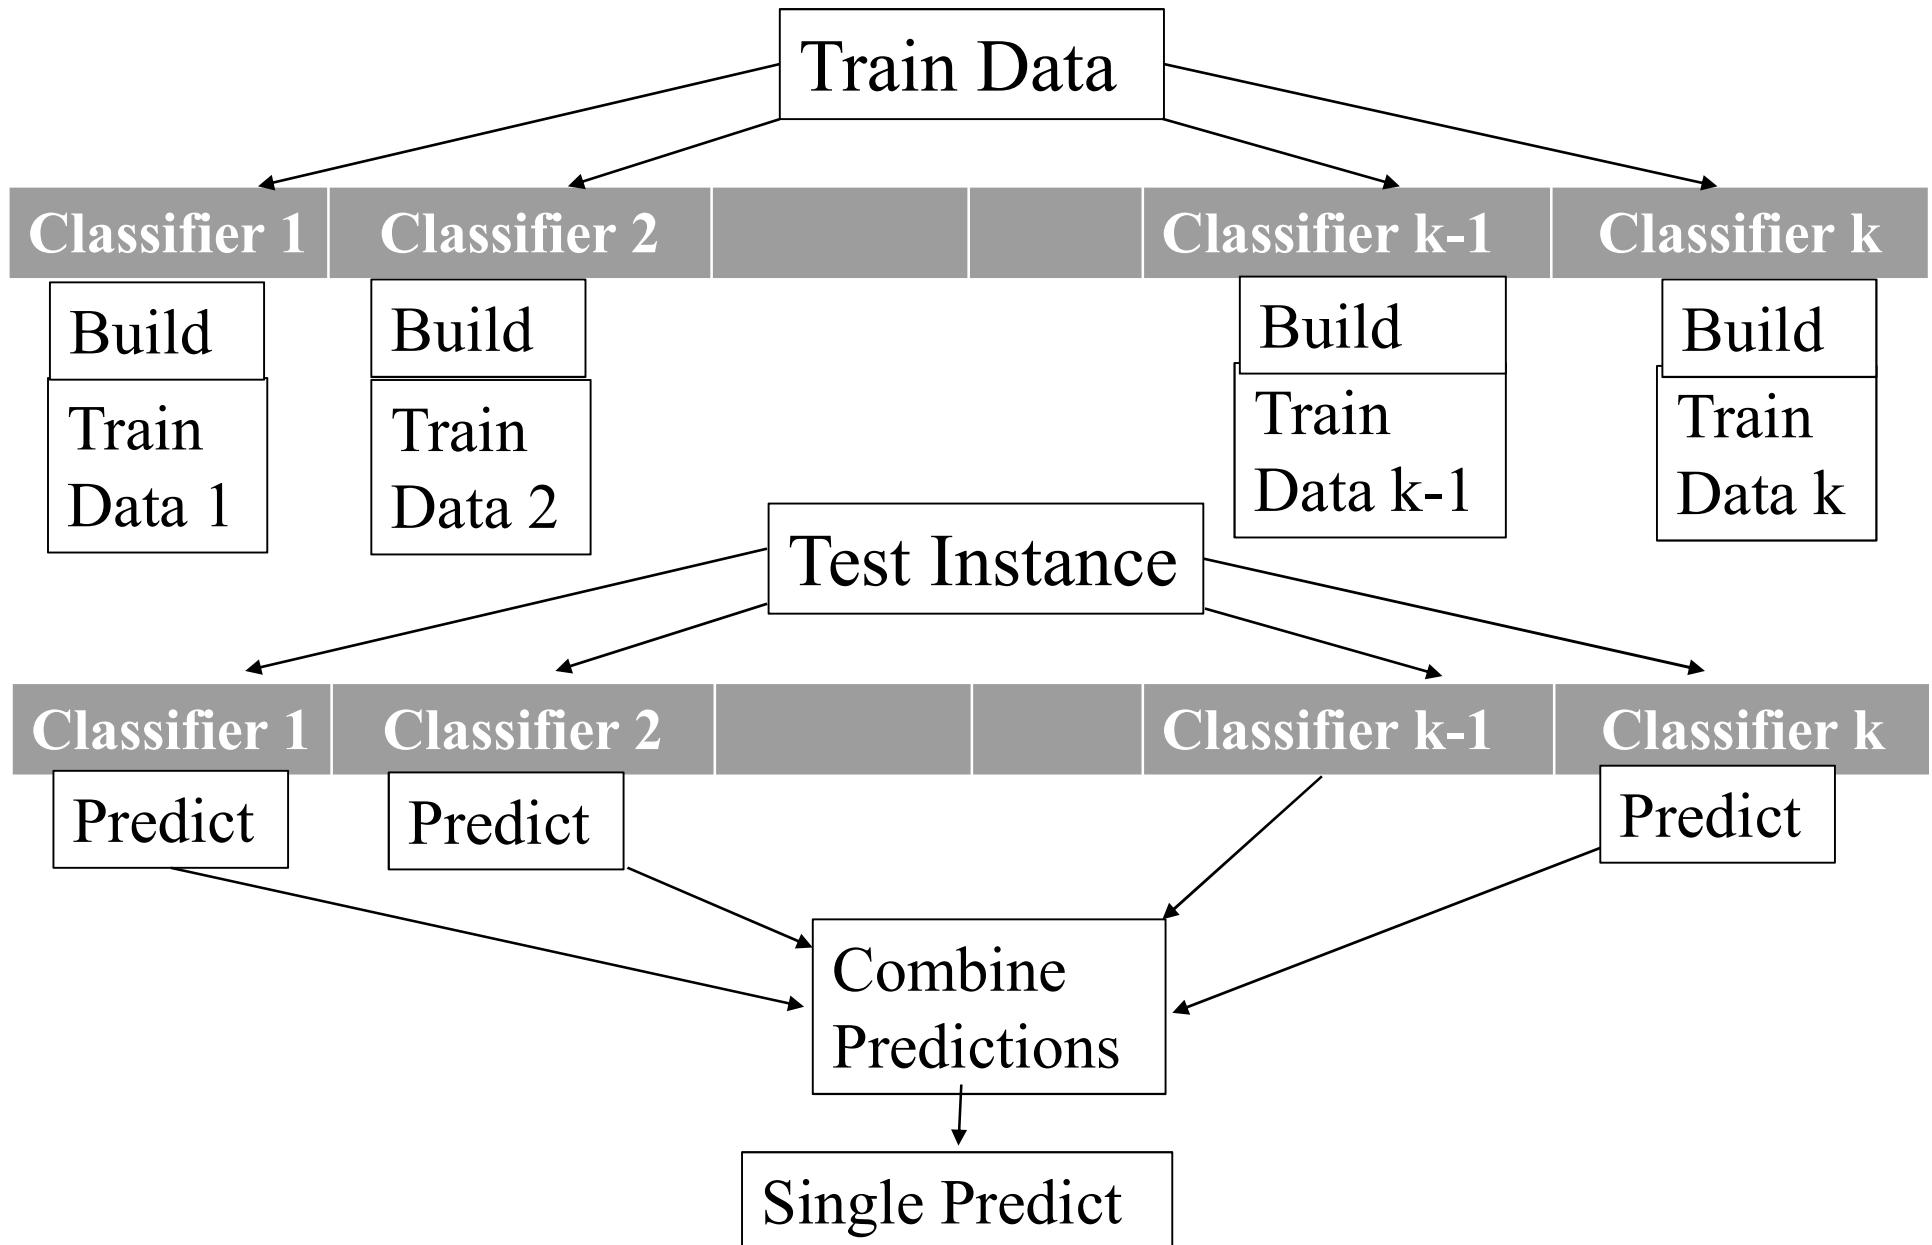
Classifier A	
Train: 0.8	
Prob Class 1	0.48
Prob Class 2	0.52
Prediction	2
Classifier B	
Train: 0.5	
Prob Class 1	0.3
Prob Class 2	0.7
Prediction	2
Classifier C	
Train: 0.75	
Prob Class 1	0.3
Prob Class 2	0.7
Prediction	2
Classifier D	
Train: 0.6	
Prob Class 1	0.7
Prob Class 2	0.3
Prediction	1
Classifier E	
Train: 0.9	
Prob Class 1	0.9
Prob Class 2	0.1
Prediction	1

Ensemble schemes are algorithms that combine the output of multiple base classifiers

Majority Vote	Class 1	Class 2
	$1 + 1 = 2$ (0.4)	$1 + 1 + 1 = 3$ (0.6)
Proportional Votes	Class 1	Class 2
	$0.6 + 0.9 = 1.5$ (0.42)	$0.8 + 0.5 + 0.75 = 2.05$ (0.58)
Weighted Probabilities ( $\alpha = 1$ )	Class 1	Class 2
	$0.8 * 0.48 + 0.5 * 0.3 + 0.75 * 0.3 + 0.6 * 0.7 + 0.9 * 0.9 = 1.99$ (0.56)	$0.8 * 0.52 + 0.5 * 0.7 + 0.75 * 0.7 + 0.6 * 0.3 + 0.9 * 0.1 = 1.56$ (0.44)
Weighted Probabilities ( $\alpha = 4$ )	Class 1	Class 2
	$0.8^4 * 0.48 + 0.5^4 * 0.3 + 0.75^4 * 0.3 + 0.6^4 * 0.7 + 0.9^4 * 0.9 = 0.99$ (0.63)	$0.8^4 * 0.52 + 0.5^4 * 0.7 + 0.75^4 * 0.7 + 0.6^4 * 0.3 + 0.9^4 * 0.1 = 0.58$ (0.37)

Ensemble Predictions		
2	X	
2	X	
1	✓	
1	✓	

# Ensemble Model



# Structural Design Choices

1. Which classifiers to use and how many?
  - Copies of the same classifier (**homogenous**)
  - Different classification algorithms  
**(heterogeneous)**
2. Which data set to train each classifier on?
  - same data/sample cases/sample attributes
3. How to combine predictions?
  - voting scheme/weighting scheme

Draw Overview Diagram  
HERE

# Which classifiers to use and how many?

There is no point combining classifiers that all produce the same predictions

Conversely, if we add too many bad classifiers, performance overall will degrade

We seek **diversity** without compromising **quality**



# Homogeneous ensembles

Bagging, random subspace and boosting

Random Forest, Rotation Forest, AdaBoost and  
XGBoost

# Diversity through sampling

## 1. Bootstrap aggregation, or bagging<sup>1</sup>

- For each classifier, sample the training data with replacement (bootstrap sampling).
- This means there will be duplicates in each training data.
- Given a training set of size  $n$ , if we bootstrap sample  $n$  times, each new training set will have approximately 63% of the cases

## 2. Random subspace methods<sup>2</sup>, or attribute sampling

- For each classifier, select a proportion of attributes to include

### Random Forest and Rotation Forest

1. L. Breiman, "Bagging predictors". *Machine Learning*. **24** (2): 123–140, 1994

2. T. K. Ho, "The Random Subspace Method for Constructing Decision Forests". *IEEE Transactions on Pattern Analysis and Machine Intelligence*. **20** (8): 832–844, 1998

# Random Forest<sub>3</sub>

- Random Forest is probably the most popular machine learning classification algorithm (after deep learning)
- It combines bagging and random subspace sampling for each base classifier
- It uses a **RandomTree** decision tree base classifier.

Stopping Criteria

Stops if a node is pure or max depth is reached, or no selected attribute gives an information gain

How to Split the Data

Different node for discrete attributes, binary split for continuous

Choosing an Attribute

**Randomly selects at most  $k$  attributes until one is found with positive information gain (not gain ratio)**

# Random Forest Details

Random Forest has the following parameters

1. Bag percent size (number of samples for each): defaults to 100
2. Number of trees: usually defaults to 500
3. Number of features: usually defaults to  $\sqrt{\text{total number of features}}$
4. Maximum tree depth: usually defaults to number of features

One benefit of Random Forests is we can assess the quality of the classifier cheaply using the **out of bag error**

Train Data 1	Train Data 2	C3	C4	Train Data k-1	Train Data k
Classifier 1	Classifier 2			Classifier k-1	Classifier k
Cases out of bag 1	Cases out of bag 2			Cases out of bag k-1	Cases out of bag k
*	1	1	*	*	0
1	*	*	*	1	*
*	*	0	*	0	*
*	*	1	*	*	1
0	1	*	0	*	*

Use C1 to predict these

The diagram shows arrows originating from the 'Cases out of bag 1' row in the main grid and pointing to the 'C1' column in the 'OOB Pred' table below. Specifically, the first three entries in the 'Cases out of bag 1' row (which are all marked with an asterisk) correspond to the first three entries in the 'C1' column of the 'OOB Pred' table (which are 1, \*, and 0 respectively). This indicates that the first three cases in the 'Cases out of bag 1' row were not included in the bootstrap sample for the first classifier (C1), and therefore their predictions are used for out-of-bag estimation.

```
package weka.classifiers.trees;
```

# Random Forest in Weka

```
/** Number of trees in forest. */
protected int m_numTrees = 10; ← Very small default number of trees

/** Number of features to consider in random feature selection.
    If less than 1 will use int(logM+1) */
protected int m_numFeatures = 0; ← This is worse than  $\sqrt{m}$ 

/** The random seed. */
protected int m_randomSeed = 1;

/** Final number of features that were considered in last build. */
protected int m_KValue = 0; ← // set up the random tree options
                           m_KValue = m_numFeatures;

/** The bagger. */
protected Bagging m_bagger = null;

/** The maximum depth of the trees (0 = unlimited) */
protected int m_MaxDepth = 0; ← Normal to leave as 0

/** The number of threads to have executing at any one time */
protected int m_numExecutionSlots = 1; ← You can thread it
```

# Random Forest in scikit

```
class RandomForestClassifier(ForestClassifier):
    """A random forest classifier.

    A random forest is a meta estimator that fits a number of decision tree
    classifiers on various sub-samples of the dataset and uses averaging to
    improve the predictive accuracy and control over-fitting.
    The sub-sample size is always the same as the original
    input sample size but the samples are drawn with replacement if
    'bootstrap=True' (default).

    def __init__(self,
                 n_estimators='warn',
                 criterion="gini",
                 max_depth=None,
                 min_samples_split=2,
                 min_samples_leaf=1,
                 min_weight_fraction_leaf=0.,
                 max_features="auto",
                 max_leaf_nodes=None,
                 min_impurity_decrease=0.,
                 min_impurity_split=None,
                 bootstrap=True,
                 oob_score=False,
                 n_jobs=None,
                 random_state=None,
                 verbose=0,
                 warm_start=False,
                 class_weight=None):
```

```
@abstractmethod
def __init__(self,
             base_estimator,
             n_estimators=100,
```

**100 default  
number of trees**

**Uses gini by default  
(not entropy)**

**Defaults to  $\sqrt{m}$**

**You can thread it**

# Rotation Forest<sup>4</sup>

Rotation forest is similar to random forest with the following key differences

1. It performs a different form of bagging and random subspace sampling, then transforms the samples using a **principle component analysis**
2. It uses C4.5 as a base classifier

You don't need to remember it in this detail

---

**Algorithm 1** buildRotationForest(Data  $D$ )

---

**Input:**  $r$ , the number of trees,  $f$ , the number of features,  $p$ , the sample proportion

- 1: Let  $\mathbf{F} = \langle F_1 \dots F_r \rangle$  be the C4.5 trees in the forest.
  - 2: **for**  $i \leftarrow 1$  to  $r$  **do**
  - 3:     Randomly partition the original features into  $k$  subsets, each with  $f$  features, denoted  $\langle S_1 \dots S_k \rangle$ .
  - 4:     Let  $D_i$  be the train set for tree  $i$ , to the original data,  $D_i = D$ .
  - 5:     **for**  $j \leftarrow 1$  to  $k$  **do**
  - 6:         Select a non-empty subset of classes and extract only cases with those class. Each class has 0.5 probability of inclusion.
  - 7:         Draw a proportion  $p$  of cases (without replacement) of those with the selected class value
  - 8:         Perform a Principal Component Analysis (PCA) on this subset of data
  - 9:         Apply the PCA transform built on this subset to the whole train set
  - 10:         Replace the features  $S_j$  in  $D_i$  with the PCA features.
  - 11:     **end for**
  - 12:     Build C4.5 Classifier  $F_i$  on transformed data  $D_i$ .
  - 13: **end for**
- 

4. Rodríguez JJ, Kuncheva LI, Alonso CJ, "Rotation forest: A new classifier ensemble method". *IEEE Transactions on Pattern Analysis and Machine Intelligence*. **28** (10): 1619–1630, 2006

# Diversity through sampling

## 3. Boosting<sup>5</sup>

- Train the first classifier on the full training data.
- Find out which cases the classifier got wrong
- For the next classifier, somehow give the wrong cases greater importance, i.e. **boost them**
- The boosting can involve resampling with increased probability, or it may be possible to embed the weights in the learning algorithm
- Boosting can also involve weighting the classifier voting
- There are many variants, these are the most popular

**AdaBoost, LogitBoost and XGBoost**

5. Schapire, Robert E. (1990). "The Strength of Weak Learnability". *Machine Learning*. 5 (2): 197–227

# Basic Boosting

---

**Algorithm 3** buildBooster(Data  $D = \langle (\mathbf{x}_1, y_1) \dots (\mathbf{x}_n, y_n) \rangle$ )

---

**Require:**  $k$ , the number of classifiers

- 1: Let  $\mathbf{H} \leftarrow \langle h_1 \dots h_k \rangle$  be the base classifiers
  - 2: Let  $\mathbf{W} \leftarrow \langle w_1 \dots w_n \rangle$  be the weights
  - 3: **for**  $i \leftarrow 1$  to  $n$  **do**
  - 4:      $w_i \leftarrow 1/n$
  - 5: **for**  $i \leftarrow 1$  to  $k$  **do**
  - 6:     buildClassifier( $D, h_i, \mathbf{W}$ )
  - 7:     updateWeights( $D, h_i, \mathbf{W}$ )
- 

- The principle is that the classifier can use the weights to alter the model it builds and that the weights are updated to give greater weight to cases the previous classifier misclassified
- One way of doing this is using the weights to sample the data

# Adaptive Boosting: AdaBoost<sup>6</sup>

---

**Algorithm 4** buildAdaBoost(Data  $D = \langle (\mathbf{x}_1, y_1) \dots (\mathbf{x}_n, y_n) \rangle$ )

---

**Require:**  $k$ , the number of classifiers

```

1: Let  $\mathbf{H} \leftarrow \langle h_1 \dots h_k \rangle$  be the base classifiers
2: Let  $\mathbf{W} \leftarrow \langle w_1 \dots w_n \rangle$  be the weights
3: for  $i \leftarrow 1$  to  $n$  do
4:    $w_i \leftarrow 1/n$ 
5: for  $i \leftarrow 1$  to  $k$  do
6:   buildClassifier( $D, h_i, \mathbf{W}$ )
7:   Let  $\mathbf{p} \leftarrow \langle p_1 \dots p_n \rangle$  be the predicted classes for the train data for  $h_i$ .
8:    $e_i \leftarrow \text{trainError}(\mathbf{y}, \mathbf{p})$ 
9:    $\alpha_i \leftarrow \frac{1}{2} \log \frac{1-e_i}{e_i}$ 
10:  for  $i \leftarrow 1$  to  $n$  do
11:     $w_i = w_i \times \begin{cases} e^{-\alpha_i} & \text{if } p_i = y_i \\ e^{\alpha_i}, & \text{otherwise} \end{cases}$ 
12:   $s_0 = 0$ 
13:  for  $i \leftarrow 1$  to  $n$  do
14:     $s_0 = s_0 + w_i$ 
15:   $s_0 = s_0/n$ 
16:  for  $i \leftarrow 1$  to  $n$  do
17:     $w_i = \frac{w_i}{s}$ 

```

The diagram illustrates the flow of the AdaBoost algorithm through five boxes connected by arrows:

- Initialise weights equally**: Box containing step 1.
- Build base classifier using weights and get predictions on train data**: Box containing steps 2-7.
- Get the overall classifier weight**: Box containing step 9.
- Tweak case weight down if correctly classified or up if misclassified**: Box containing steps 11-13.
- Normalise so weights add up to one**: Box containing steps 15-17.

---

6. Freund and Schapire: Experiments with a new boosting algorithm. In: Thirteenth International Conference on Machine Learning, San Francisco, 148-156, 1996.

# AdaBoost Predictions

New cases are classified by a **weighted majority vote**

---

**Algorithm 5** classifyInstanceAdaBoost(Data  $\mathbf{x}$ )

---

```
1:  $score \leftarrow <0, 0, \dots, 0>$  //Number of classes  
2: for  $h_i \in \mathbf{H}$  do  
3:    $p \leftarrow \text{classifyInstance}(h_i, \mathbf{x})$   
4:    $score_p \leftarrow score_p + \alpha_i$   
return  $\arg \max(score)$ 
```

---

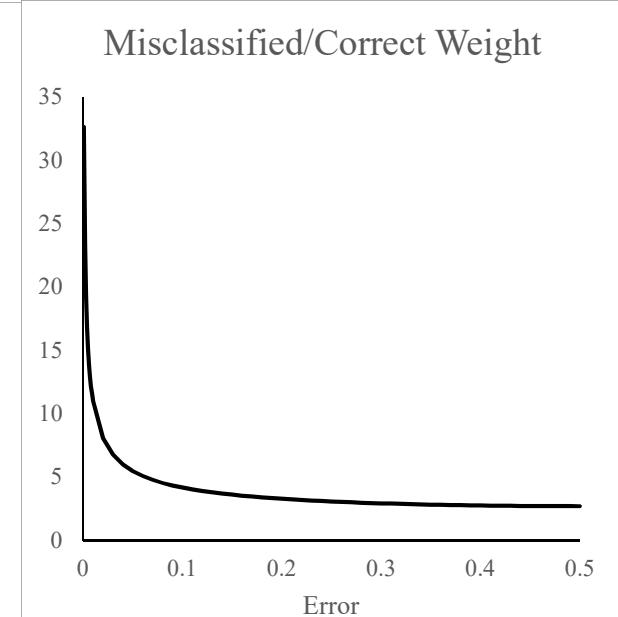
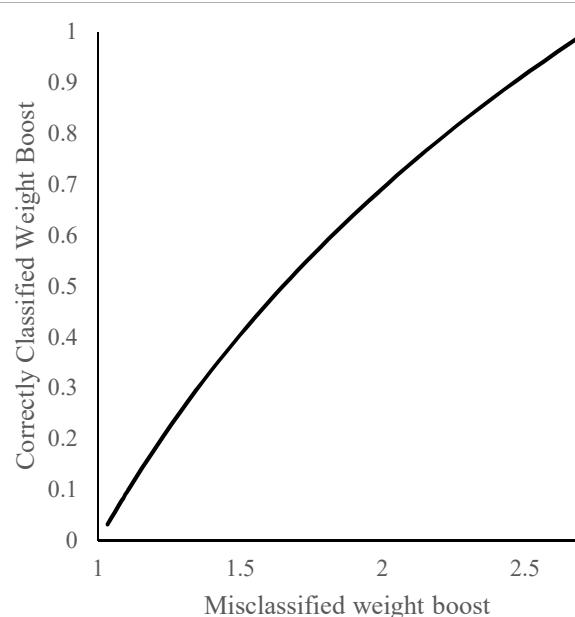
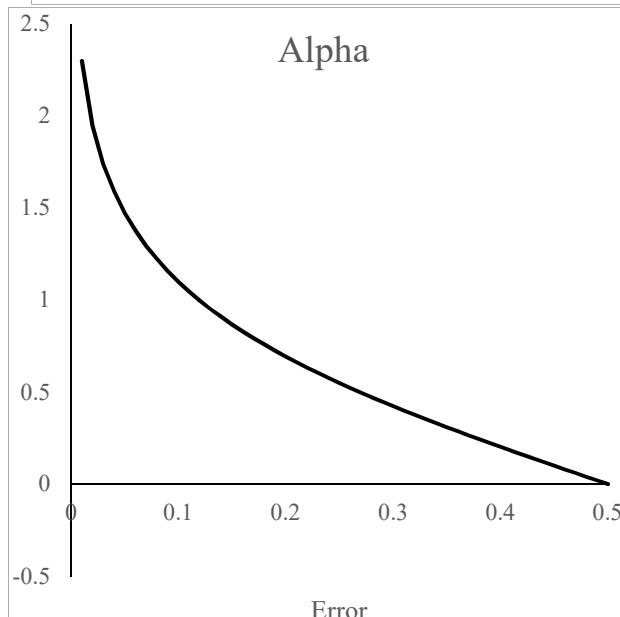
- Each ensemble member is given weight equal to alpha.

Weights can be used by the base classifier in two ways

1. Internally, to adjust the way the classifier builds the model
2. Externally, by using them to sample through bagging

# AdaBoost Details

- The base classifier in the original work and the Weka implementation is a **Decision Stump**
- A decision stump is a 1 level classification tree that uses information gain
- Weights can be used by the base classifier in two ways
  1. Internally, to adjust the way the classifier builds the model
  2. Externally, by using them to sample through bagging (default with stump)



# Other Popular Boosting Algorithms

## Logistic Boosting: LogitBoost

- Use a decision stump like AdaBoost
- Adjusts the weights after each new model using and adaption of Logistic Regression
- Always does 1 vs All models

## Extreme Gradient Boost: XGBoost

- Currently popular because of Kaggle success
- Uses CART trees as base models
- Uses a greedy algorithm to create an additive model, each iteration tries to correct the errors of the previous
- Uses a penalty function (**regularisation**) to avoid overfitting the train data

# Heterogeneous ensembles

**Weighting Schemes:** majority vote, weighted majority vote, weighted probabilities (CAWPE)

**Selection Schemes:** overproduce and select

**Stacking:** build a model on the output of the base models

Ignore the base classifiers, address the issue of how to combine the classifier outputs to improve performance

**Assume we have  $k$  classifiers**

# Weighting Schemes

Assume  $k$  base classifiers. For a new case they produce  $k$  predictions and  $k$  probability distributions

Classifier A	
Train: 0.8	
Prob Class 1	0.48
Prob Class 2	0.52
Prediction	2
Classifier B	
Train: 0.5	
Prob Class 1	0.3
Prob Class 2	0.7
Prediction	2
Classifier C	
Train: 0.75	
Prob Class 1	0.3
Prob Class 2	0.7
Prediction	2
Classifier D	
Train: 0.6	
Prob Class 1	0.7
Prob Class 2	0.3
Prediction	1
Classifier E	
Train: 0.9	
Prob Class 1	0.9
Prob Class 2	0.1
Prediction	1

**Majority voting** simply counts how many vote for each class

Majority Vote →

Class 1	Class 2
$1 + 1 =$ 2 (0.4)	$1 + 1 + 1 =$ 3 <u>(0.6)</u>

**Weighted majority voting** gives each classifier a weight formed on the train data (we estimate accuracy with 10-fold cross validation)

Proportional Votes →

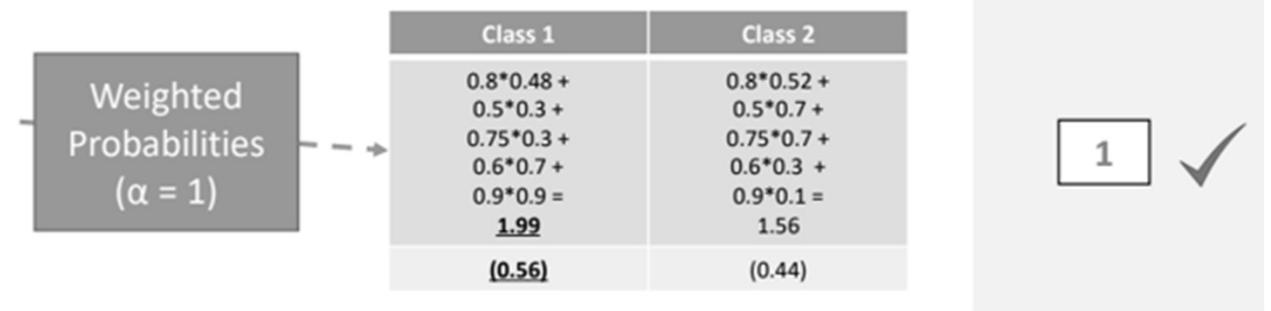
Class 1	Class 2
$0.6 + 0.9 =$ 1.5 (0.42)	$0.8 + 0.5 + 0.75 =$ 2.05 <u>(0.58)</u>

# Weighting Schemes

**Weights based on an estimate of the test accuracy found on the train data through cross validation or bagging**

Classifier A	
Train: 0.8	
Prob Class 1	0.48
Prob Class 2	0.52
Prediction	2
Classifier B	
Train: 0.5	
Prob Class 1	0.3
Prob Class 2	0.7
Prediction	2
Classifier C	
Train: 0.75	
Prob Class 1	0.3
Prob Class 2	0.7
Prediction	2
Classifier D	
Train: 0.6	
Prob Class 1	0.7
Prob Class 2	0.3
Prediction	1
Classifier E	
Train: 0.9	
Prob Class 1	0.9
Prob Class 2	0.1
Prediction	1

**Probability weighting** applies the weights to the probabilities, not the predictions



**CAWPE** raises the weights to the power alpha. It uses a tilted distribution and inverse temperature



# Selection Schemes

- A currently popular approach is to build a large set of classifiers, then select a subset to actually use.
- This is often called an “overproduce and choose” scheme.
- Often they use a greedy algorithm to construct a model

Classifier 1	Classifier 2			Classifier T-1	Classifier T
Acc = 0.87	Acc=0.67			Acc=0.60	Acc=0.91

Classifier T

Classifier T

Classifier 8

Classifier T

Classifier 8

Classifier 3

# Stacking

- Stacking involves constructing a “meta” model on top of the training data output.
- We perform a cross validation on the train data, which gives us predictions and probabilities for each train case
- Suppose we have three base classifiers and we cross validate to get predictions and probabilities on the train data

Random Forest

0	1	0.166	0.834
0	0	0.802	0.198
0	0	0.638	0.362
0	0	0.54	0.46
1	1	0.096	0.904
1	0	0.852	0.148
1	0	0.574	0.426

MLP

0	1	0.021619	0.978381
0	0	0.995253	0.004747
0	0	0.778656	0.221344
0	1	0.073482	0.926518
1	1	0.036168	0.963832
1	0	0.990263	0.009737
1	0	0.971657	0.028343

C4.5

0	1	0	1
0	0	0.666667	0.333333
0	1	0.25	0.75
0	1	0.25	0.75
1	0	0.666667	0.333333
1	0	1	0
1	0	1	0

# Stacking

- We use the output to form a new classification problem.  
Assume we use the probabilities

0.166	0.834	0.022	0.978	0	1	0
0.802	0.198	0.995	0.005	0.667	0.333	0
0.638	0.362	0.779	0.221	0.25	0.75	0
0.54	0.46	0.073	0.927	0.25	0.75	0
0.096	0.904	0.036	0.964	0.667	0.333	1
0.852	0.148	0.99	0.01	1	0	1
0.574	0.426	0.972	0.028	1	0	1



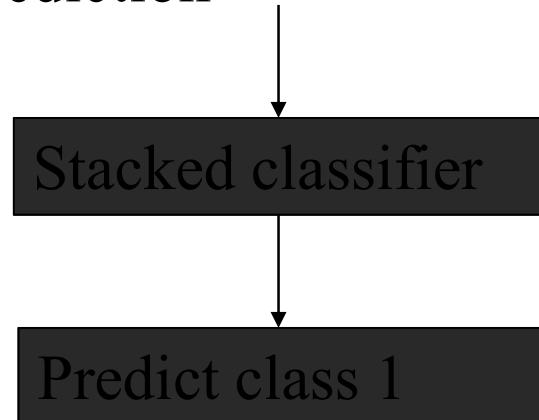
Train stacked classifier

# Stacking

- New cases are classified by first finding the output of the base classifiers

0.11    0.89    0.44    0.56    0.76    0.24

- This output is passed to the “meta” classifier, which then produces a prediction



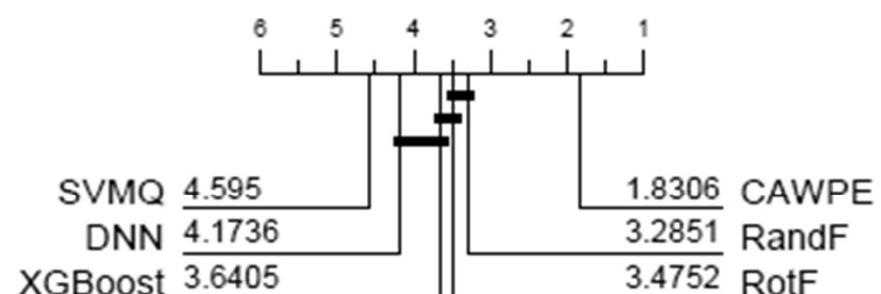
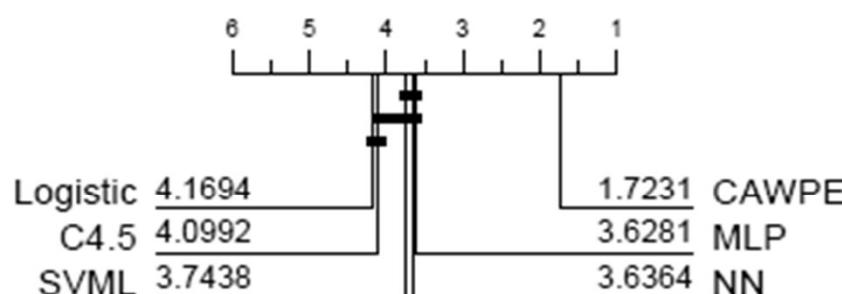
---

**Algorithm 1** CAWPE classify(A test case  $x$ )

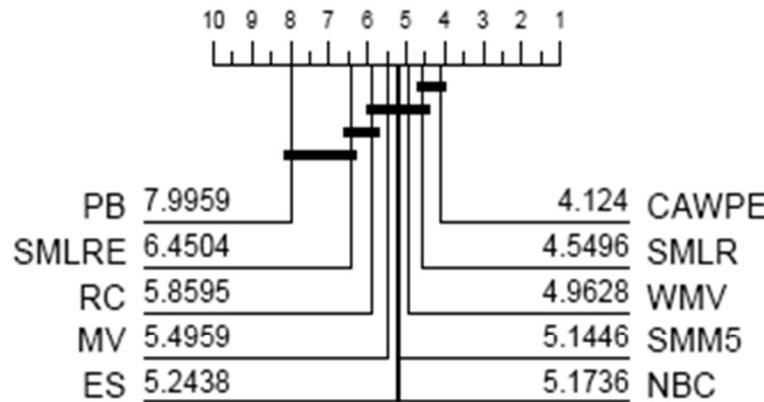
---

**Output:** prediction for case  $x$ )

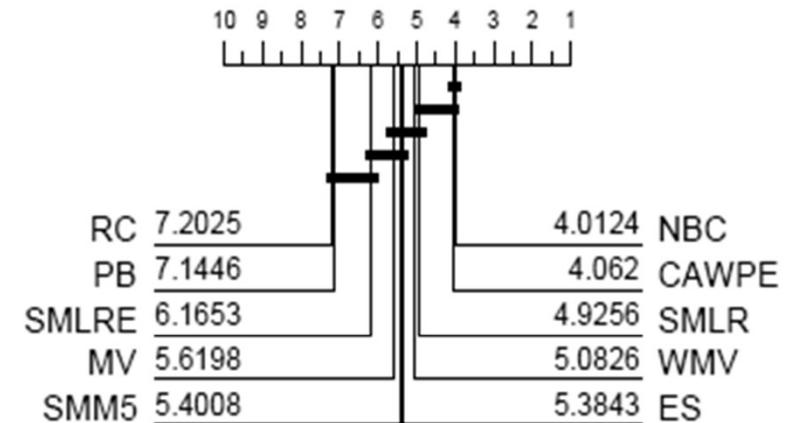
- 1: Given a set of classifiers  $\langle M_1, \dots, M_k \rangle$ , an exponent  $\alpha$ , a set of weights,  $w_i$  and the number of classes  $c$
  - 2:  $\{\hat{p}_1, \dots, \hat{p}_c\} = \{0, \dots, 0\}$
  - 3: **for**  $i \leftarrow 1$  to  $k$  **do**
  - 4:     **for**  $j \leftarrow 1$  to  $c$  **do**
  - 5:          $\hat{q}_j \leftarrow \hat{p}((y=j|M_i, x)$
  - 6:          $\hat{p}_j \leftarrow \hat{p}_j + w_i^\alpha \cdot \hat{q}_j$
- return**  $\arg \max_{j=1 \dots c} \hat{p}_j$
- 



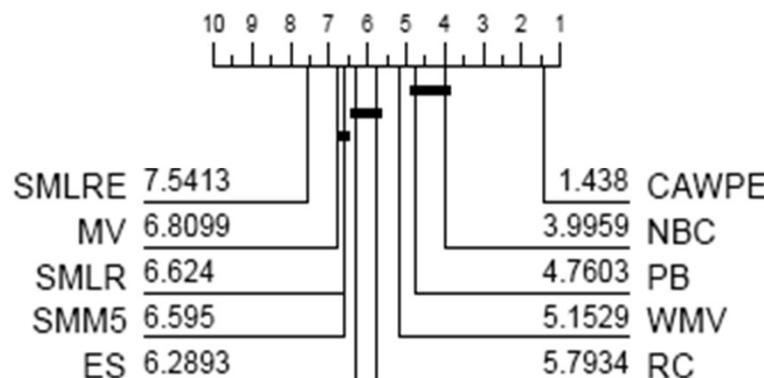
# Comparison of Weighting Schemes for Heterogeneous Classifiers



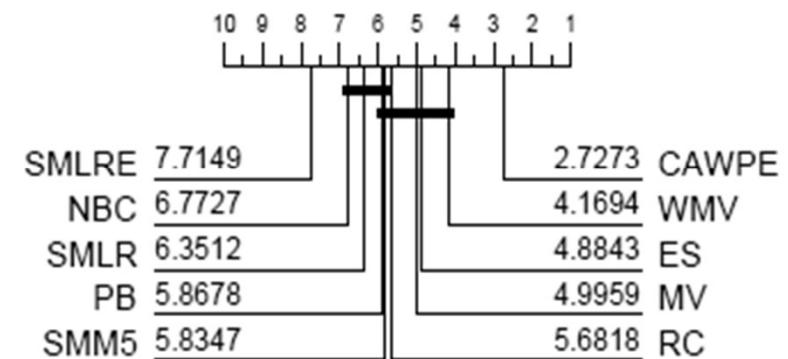
(a) Error



(b) Balanced Error

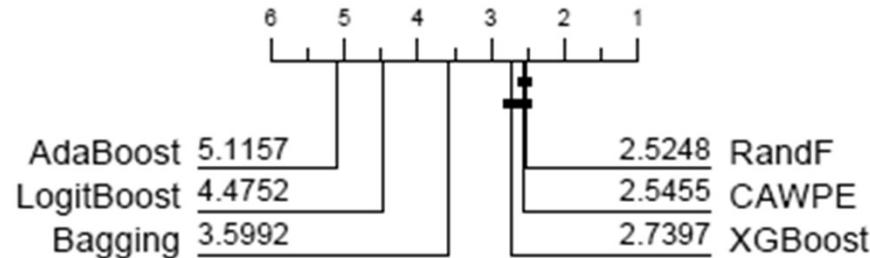


(c) AUC

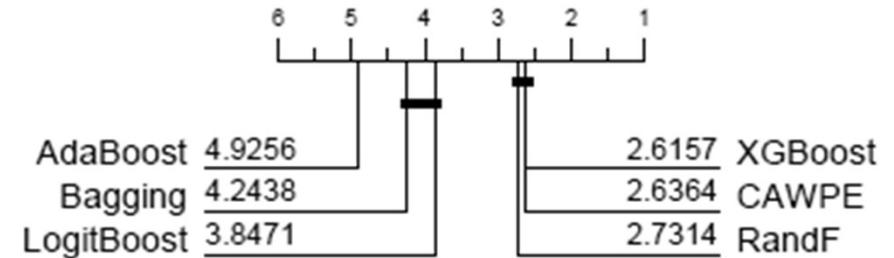


(d) NLL

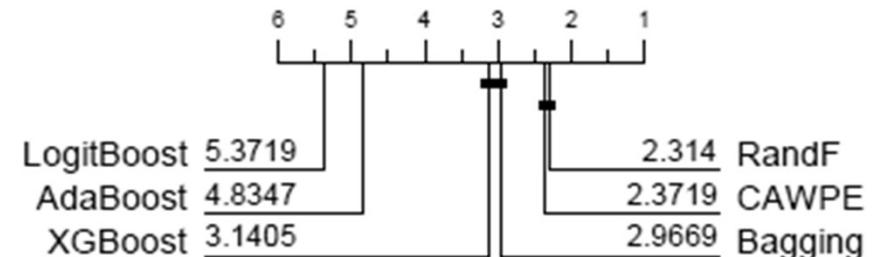
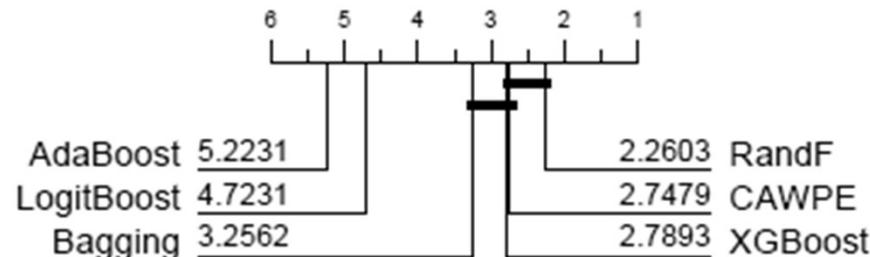
# Comparison of Homogeneous classifiers (and basic CAWPE)



(a) Error



(b) Balanced Error



# Comparison of tuned classifiers

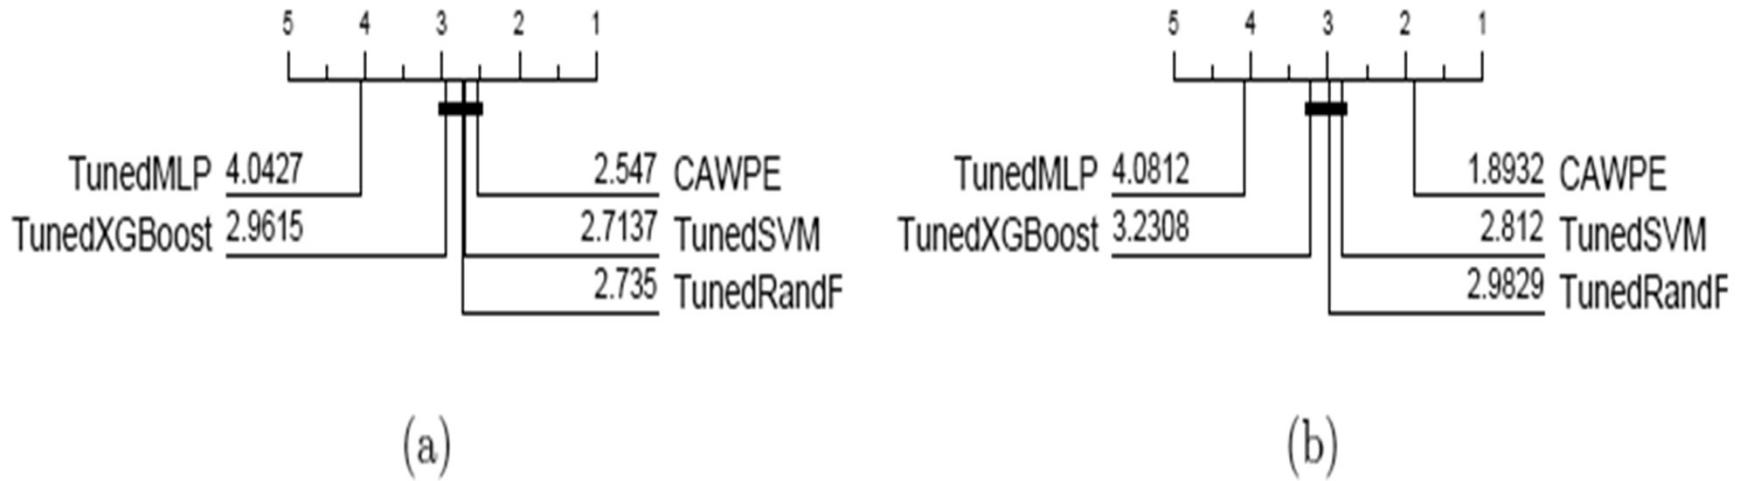


Fig. 6 Average ranked errors for CAWPE with (a) simpler and (b) advanced components against four tuned classifiers on 117 datasets in the UCI archive. The datasets adult, chess-krvk, miniboone and magic are omitted due to computational restraints.

# Ensembles Summary

**The design goal is diversity without loss of accuracy**

**Ensembles main purpose is to reduce the variability in estimates, but they are usually more accurate too.**

**Homogeneous** ensembles use a large number of “weak” learners and inject diversity through sampling the data.

*Samplers:* *Random Forest, Rotation Forest*

*Boosters:* *AdaBoost, LogitBoost, XGBoost*

**Heterogeneous** classifiers use a smaller number of “strong” classifiers

*Weighting Scheme: CAWPE*

*Overproduce and Select*

*Stacking*