

CMP-6002B - Machine Learning

Dr Gavin Cawley

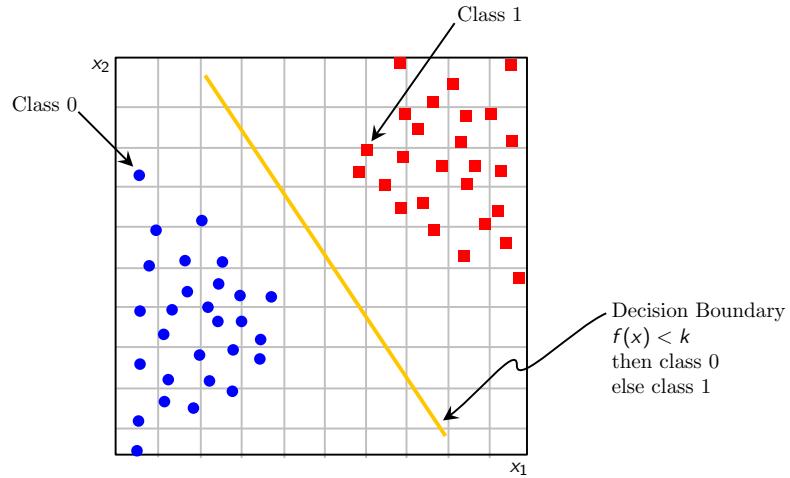
Lecture 4 - Linear Classifiers

Introduction

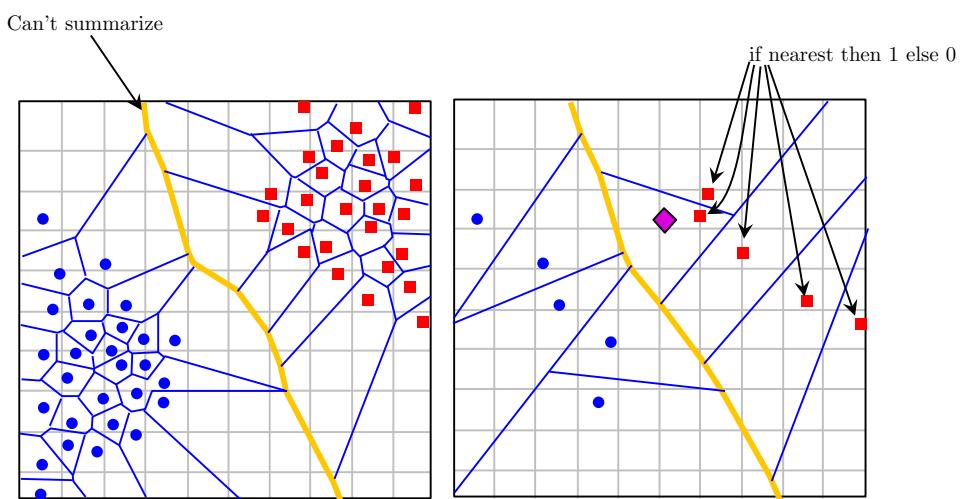
- ▶ Objective:
 - ▶ To understand the principle and practice of linear classifiers.
- ▶ Desired Outcomes:
 - ▶ Understand the principles of linear discriminant analysis
 - ▶ Understand the principles and practice of the perceptron approach
- ▶ Overview of Linear Classifiers
 - ▶ Optimal classifiers for Gaussian data
 - ▶ Fisher's linear discriminant
 - ▶ Linear classifiers as regression
 - ▶ Linear classifiers as Perceptrons
- ▶ Strengths and weaknesses
- ▶ Extensions to non-linear classifiers

Decision Boundaries

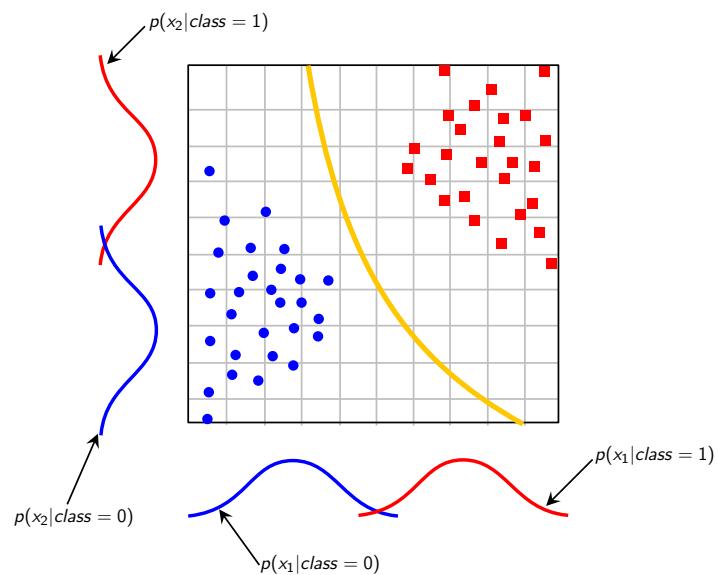
- ▶ The general problem of classification is to most effectively separate the classes.



1-Nearest Neighbour Classifier



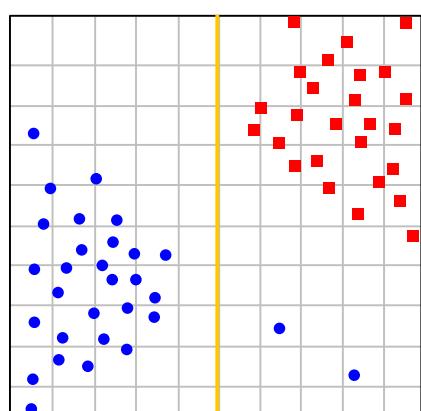
Naïve Bayes Classifier



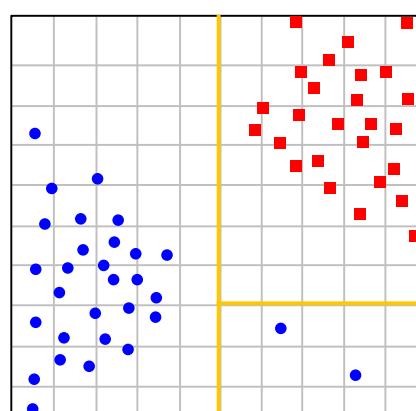
$$p(class = 0|x_1, x_2) = \frac{p(x_1|class = 0) \times p(x_2|class = 0) \times p(class = 0)}{p(x_1) \times p(x_2)}$$

Decision Tree Classifier

```
if ( $x_1 < 5$ ) then  
    Circle  
else  
    Square
```



```
if ( $x_1 < 5$ ) then  
    Circle  
else  
    if ( $x_2 > 3$ ) then  
        Square  
    else  
        Circle
```

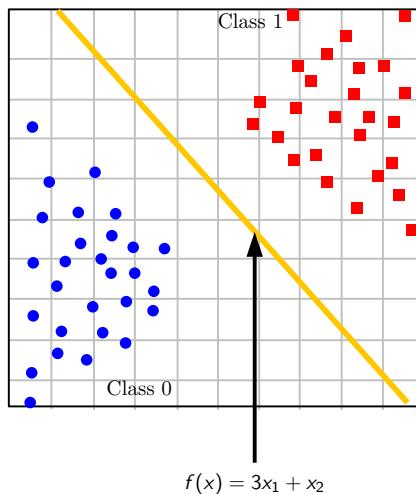


Linear Classifiers

- ▶ A linear model takes the form

$$g(\mathbf{x}) = f(\mathbf{w} \cdot \mathbf{x}) = f\left(\sum_{i=1}^d w_i x_i\right)$$

where \mathbf{w} are the model parameters



Discriminant Functions

- ▶ A classifier can be represented using *discriminant* functions

$$g_1(\mathbf{x}), g_2(\mathbf{x}), \dots, g_c(\mathbf{x}).$$

- ▶ Assign to the class with the most positive discriminant,

$$g_i(\mathbf{x}) > g_j(\mathbf{x}) \quad \forall i \neq j$$

- ▶ For a probabilistic classifier, the error rate is minimised if

$$g_i(\mathbf{x}) = p(\mathcal{C}_i | \mathbf{x}) = p(\mathbf{x} | \mathcal{C}_i) p(\mathcal{C}_i) / p(\mathbf{x})$$

- ▶ Not affected by monotonic transformations or constants, so

$$g_i(\mathbf{x}) = \log p(\mathbf{x} | \mathcal{C}_i) + \log p(\mathcal{C}_i)$$

- ▶ Single discriminant for 2-class problems,

$$g(\mathbf{x}) = g_1(\mathbf{x}) - g_2(\mathbf{x}).$$

The Normal or Gaussian Density

- ▶ Linear classifiers can be motivated using Gaussian densities.
- ▶ Gaussian densities often used for mathematical tractability.
- ▶ Univariate Gaussian; parameters: mean μ and variance σ^2

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{(x - \mu)^2}{2\sigma^2} \right\}$$

- ▶ Multivariate Gaussian,

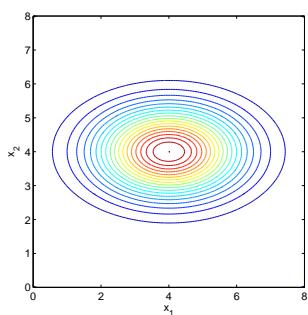
$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2}|\boldsymbol{\Sigma}|^{1/2}} \exp \left\{ -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

Parameters:

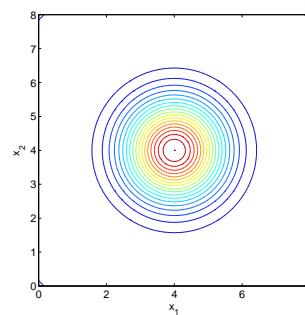
- $\boldsymbol{\mu}$ - Mean vector
 $\boldsymbol{\Sigma}$ - Covariance matrix

Multivariate Gaussian Densities

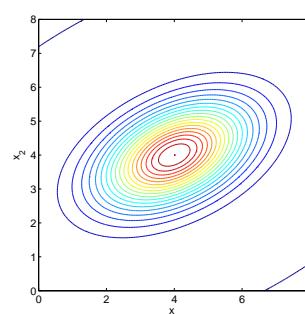
$$\begin{aligned} \boldsymbol{\mu} &= [4 \ 4] \\ \boldsymbol{\Sigma} &= \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \end{aligned}$$



$$\begin{aligned} \boldsymbol{\mu} &= [4 \ 4] \\ \boldsymbol{\Sigma} &= \begin{bmatrix} 2 & 0 \\ 0.5 & 1 \end{bmatrix} \end{aligned}$$



$$\begin{aligned} \boldsymbol{\mu} &= [4 \ 4] \\ \boldsymbol{\Sigma} &= \begin{bmatrix} 2 & 0 \\ 0 & 0.75 \end{bmatrix} \end{aligned}$$



Identical Spherical Gaussian Densities

- ▶ Assume the class conditional densities are spherical Gaussians with equal variances

$$p(\mathbf{x}|\mathcal{C}_1) \sim \mathcal{N}(\boldsymbol{\mu}_1, \sigma^2 \mathbf{I}) \quad \text{and} \quad p(\mathbf{x}|\mathcal{C}_2) \sim \mathcal{N}(\boldsymbol{\mu}_2, \sigma^2 \mathbf{I})$$

- ▶ Then the optimal discriminant functions are given by

$$g_i(\mathbf{x}) = -\|\mathbf{x} - \boldsymbol{\mu}_i\|^2 / (2\sigma^2) + \log p(\mathcal{C}_i)$$

- ▶ Noting that $\|\mathbf{x} - \boldsymbol{\mu}_i\|^2 = (\mathbf{x} - \boldsymbol{\mu}_i)^T (\mathbf{x} - \boldsymbol{\mu}_i)$,

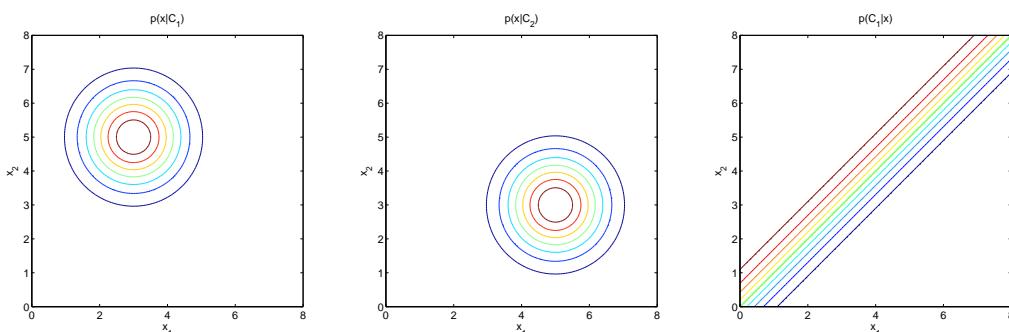
$$g_i(\mathbf{x}) = -[\mathbf{x}^T \mathbf{x} - 2\boldsymbol{\mu}_i^T \mathbf{x} + \boldsymbol{\mu}_i^T \boldsymbol{\mu}_i] / (2\sigma^2) + \log p(\mathcal{C}_i)$$

- ▶ So a linear discriminant is optimal, $g_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + b_i$, where

$$\mathbf{w}_i = \boldsymbol{\mu}_i / \sigma^2 \quad \text{and} \quad b_i = -\boldsymbol{\mu}_i^T \boldsymbol{\mu}_i / (2\sigma^2) + \log p(\mathcal{C}_i)$$

Example: Identical Spherical Gaussian Densities

- ▶ For this example: $\boldsymbol{\mu}_1 = [3 \ 5]$, $\boldsymbol{\mu}_2 = [5 \ 3]$ and $\boldsymbol{\Sigma}_1 = \boldsymbol{\Sigma}_2 = \mathbf{I}$.
- ▶ Equal prior probabilities $p(\mathcal{C}_1) = p(\mathcal{C}_2) = 0.5$.



- ▶ The contours of the posterior $p(\mathcal{C}_1|\mathbf{x})$ are linear
- ▶ The optimal discriminant function is also linear
 - ▶ $g(\mathbf{x}) = 0$ corresponds to $p(\mathcal{C}_1|\mathbf{x}) = 0.5$.

Example: Identical Spherical Gaussian Densities

- ▶ Step 1: construct $g_1(\mathbf{x}) = \mathbf{w}_1^T \mathbf{x} + b_1$.

$$\mathbf{w}_1 = \boldsymbol{\mu}_1/\sigma^2 = [3 \ 5], \quad b_1 = -\boldsymbol{\mu}_1^T \boldsymbol{\mu}_1/(2\sigma^2) \approx -15.6931$$

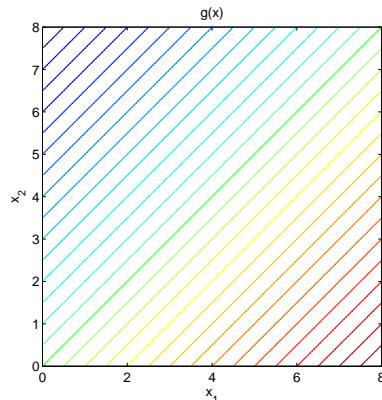
- ▶ Step 2: construct $g_2(\mathbf{x}) = \mathbf{w}_2^T \mathbf{x} + b_2$.

$$\mathbf{w}_2 = \boldsymbol{\mu}_2/\sigma^2 = [5 \ 3], \quad b_2 = -\boldsymbol{\mu}_2^T \boldsymbol{\mu}_2/(2\sigma^2) \approx -15.6931$$

- ▶ Step 3: construct single discriminant

$$\begin{aligned} g(\mathbf{x}) &= g_1(\mathbf{x}) - g_2(\mathbf{x}) \\ &= [5 \ 3]^T \mathbf{x} - [3 \ 5]^T \mathbf{x} \\ &= 2x_1 - 2x_2 \end{aligned}$$

- ▶ c.f. $p(\mathcal{C}_1|\mathbf{x}) = 0.5$.



Identical Arbitrary Gaussians

- ▶ Suppose the densities have different means, but common covariance $\boldsymbol{\Sigma}_i = \boldsymbol{\Sigma}$.
- ▶ Again we write the discriminant function,

$$g_1(\mathbf{x}) = -(\mathbf{x} - \boldsymbol{\mu}_1)^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}_1)/2 + \log p(\mathcal{C}_1).$$

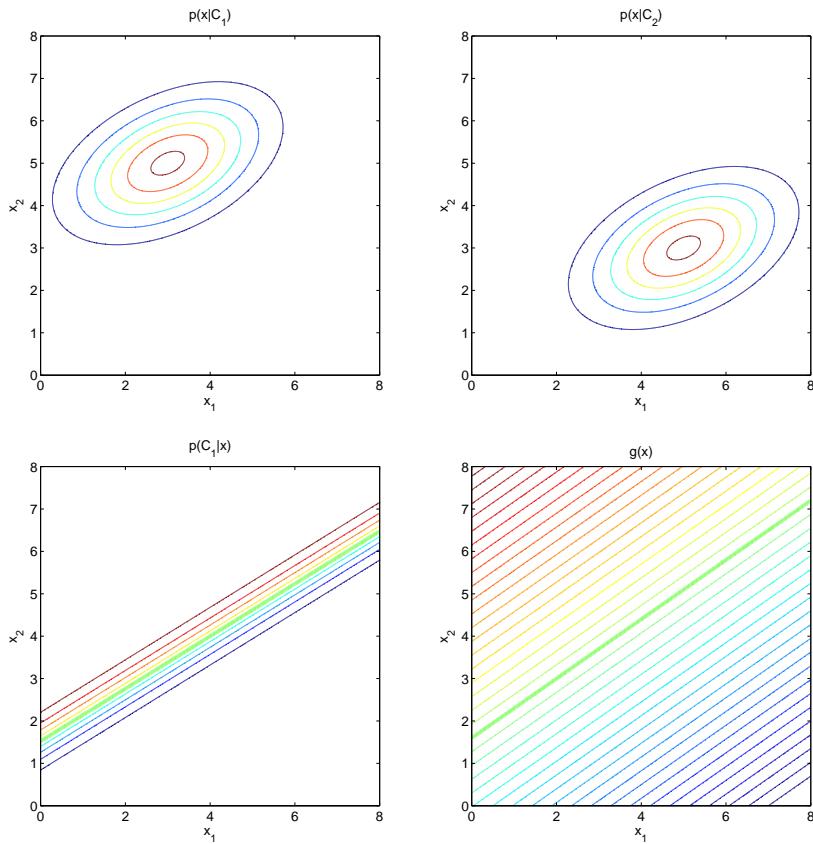
- ▶ Multiplying out, we obtain a common term $\mathbf{x}^T \boldsymbol{\Sigma}^{-1} \mathbf{x}$.
- ▶ Again a linear discriminant is optimal,

$$g_1(\mathbf{x}) = \mathbf{w}_1^T \mathbf{x} + b_1,$$

where

$$\mathbf{w}_1 = \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}_1 \quad \text{and} \quad b_1 = -\boldsymbol{\mu}_1^T \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}_1/2 + \log p(\mathcal{C}_1).$$

Example: Identical Arbitrary Gaussian Densities



Arbitrary Gaussian Densities

- ▶ Consider the case where Σ_1 and Σ_2 are arbitrary covariance matrices.
- ▶ In this case we have a quadratic classifier

$$g_i(\mathbf{x}) = \mathbf{x}^T \mathbf{W}_i \mathbf{x} + \mathbf{w}_i^T \mathbf{w}_i + b_i$$

where

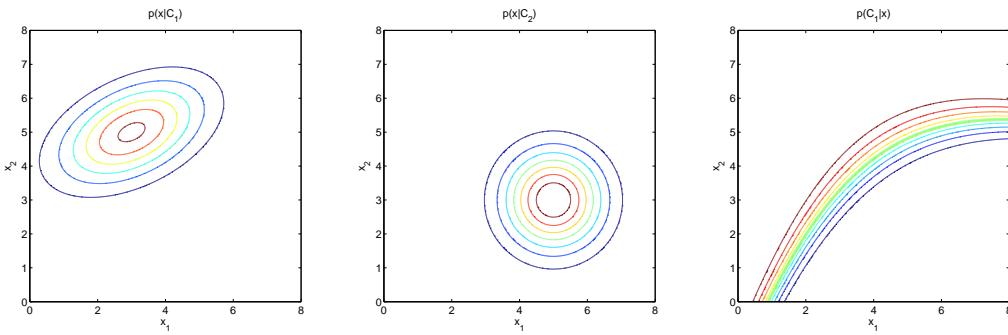
$$\mathbf{W}_i = -\frac{1}{2}\Sigma_i^{-1}, \quad \mathbf{w}_i = \Sigma_i^{-1}\mu_i,$$

and

$$b_i = -\frac{1}{2}\mu_i^T \Sigma_i^{-1} \mu_i - \frac{1}{2} \log |\Sigma_i| + \log(C_i).$$

- ▶ The optimal discriminant will be a hyper-quadratic.

Example: Arbitrary Gaussian Densities



- ▶ In two dimensions, a quadratic discriminant is of the form

$$g(\mathbf{x}) = \omega_1 x_1^2 + \omega_2 x_2^2 + \omega_3 x_1 x_2 + \omega_4 x_1 + \omega_5 x_2 + b$$

- ▶ Can implement as a linear discriminant in an augmented attribute space.

Linear Discriminant Analysis

- ▶ Fisher's Linear Discriminant Analysis¹ (LDA) is *the* classic linear classifier.
- ▶ Basic idea: Find the linear projection of the data that maximises the distance between patterns of different classes while minimising the distance between patterns of the same class.
- ▶ Consider the projection onto a line joining the class means,

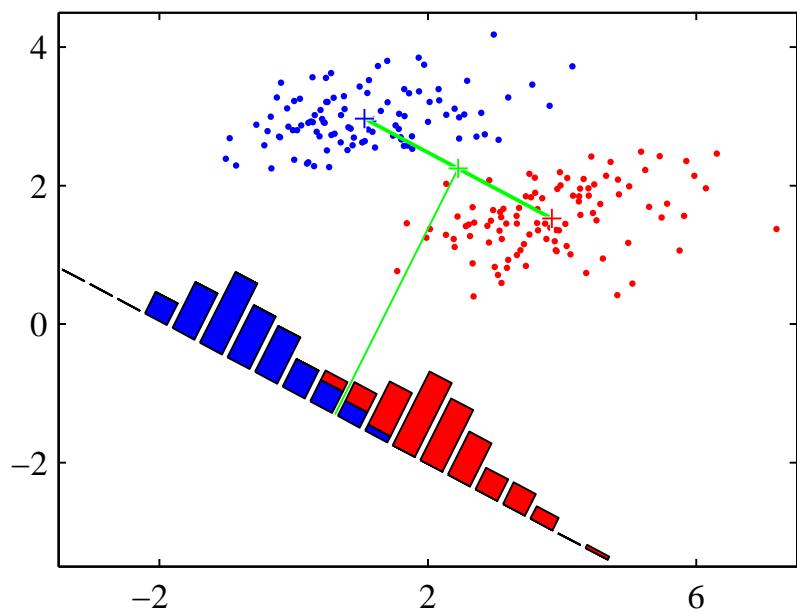
$$g(\mathbf{x}) = (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^T \mathbf{x}$$

where

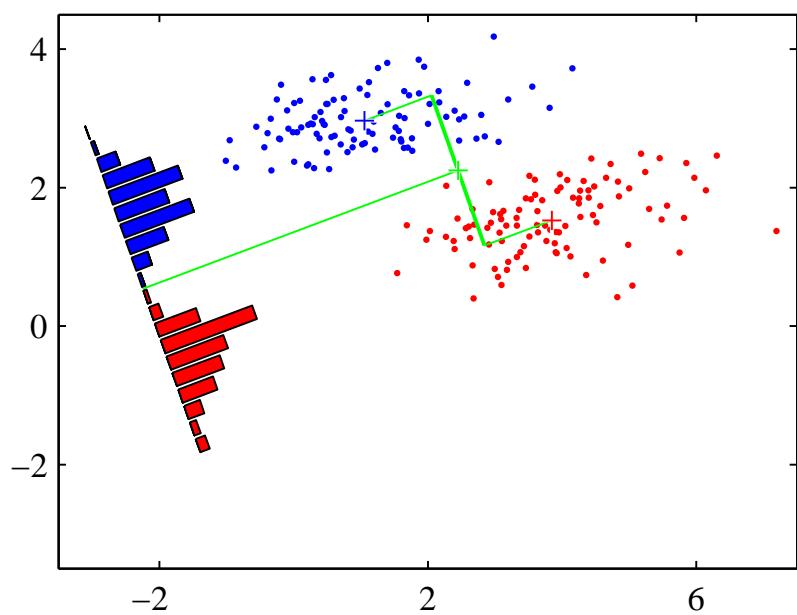
$$\boldsymbol{\mu}_1 = \frac{1}{N_1} \sum_{i \in \mathcal{C}_1} \mathbf{x}_i \quad \text{and} \quad \boldsymbol{\mu}_2 = \frac{1}{N_2} \sum_{i \in \mathcal{C}_2} \mathbf{x}_i$$

¹Fisher, R.A. The Use of Multiple Measurements in Taxonomic Problems. Annals of Eugenics, 7: 179-188 (1936)

Projection onto Line Joining Class Means



A Better Linear Discriminant



Linear Discriminant Analysis

- ▶ The second discriminant takes more account of the *distribution* of positive and negative patterns.
- ▶ Linear discriminant $g(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$, but how to choose \mathbf{w} ?
- ▶ Optimise the Fisher ratio

$$\mathcal{J}(\mathbf{w}) = \frac{(m_1 - m_2)^2}{s_1^2 + s_2^2}$$

- ▶ Let $y_n = \mathbf{w}^T \mathbf{x}_n$, then

$$m_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} y_n \quad \text{and} \quad s_k^2 = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} (y_n - m_k)^2$$

- ▶ Distance between class means divided by average distance to class mean.

Fisher's Linear Discriminant Analysis

- ▶ We can rewrite the Fisher criterion as

$$\mathcal{J}(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

- ▶ \mathbf{S}_B is the *between class scatter matrix*

$$\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

- ▶ \mathbf{S}_W is the *within class scatter matrix*

$$\mathbf{S}_W = \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^T + \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^T$$

- ▶ Differentiate Fisher criterion w.r.t. \mathbf{w} and set to zero

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w} \implies \mathbf{w} \propto \mathbf{S}_W^{-1} (\mathbf{m}_2 - \mathbf{m}_1)$$

Least Squares Linear Classifiers

- Given m attributes x_1, x_2, \dots, x_m and observed responses y , estimate parameters $w_0, w_1, w_2, \dots, w_m$ for the line

$$g(\mathbf{x}) = w_0 + w_1 x_1 + \dots + w_m x_m$$

- Fit a line, for example

- Line 1:

$$g_a(\mathbf{x}) = 0.2 + 0.3 \times x_1 - 0.2 \times x_2$$

- Line 2:

$$g_b(\mathbf{x}) = -0.3 + 0.13 \times x_1 + 0.05 \times x_2$$

- How do we choose which is best?

x_1	x_2	y
2.7	5.5	0
0.9	4.7	0
1.1	3.1	0
2.9	1.9	0
0.5	1.0	0
8.0	9.1	1
5.4	8.5	1
6.1	6.6	1
8.3	6.6	1
8.1	4.7	1

Linear Classifiers via Multiple Regression

- Evaluate predicted or fitted values

y	0	0	0	0	0	1	1	1	1	1
$g_a(\mathbf{x})$	0.06	-0.08	0.78	1.92	1.65	-0.15	-0.63	0.53	1.19	2.08
$g_b(\mathbf{x})$	0.29	0.03	-0.02	0.16	-0.19	1.14	0.78	0.78	1.07	0.96

- From the fitted values we can form a decision boundary by classifying a case as the nearest fitted integer value

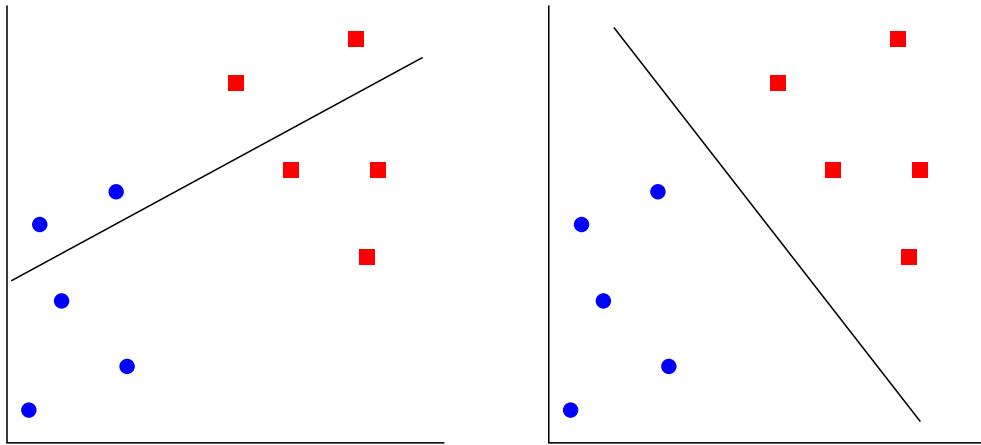
y	0	0	0	0	0	1	1	1	1	1
$g_a(\mathbf{x})$	0	0	1	1	1	0	0	1	1	1
$g_b(\mathbf{x})$	0	0	0	0	0	1	1	1	1	1

- $g_b(\mathbf{x})$ is the better discriminant.

Linear Classifiers via Multiple Regression

So, for any line, the decision boundary for a binary response variable is simply

$$w_0 + w_1 x_1 + \cdots + w_m x_m < 0.5$$



Linear Classifiers via Multiple Regression

- ▶ Which of the infinite possible lines should we choose?
- ▶ The obvious candidate is the one that *minimises the sum of the squared error*

$$SSE = \sum_{i=1}^n (y_i - g(\mathbf{x}_i))^2$$

- ▶ Compute SSE for each discriminant:

y	0	0	0	0	0	1	1	1	1	1
$g_a(\mathbf{x})$	0.06	-0.08	0.78	1.92	1.65	-0.15	-0.63	0.53	1.19	2.08
$g_b(\mathbf{x})$	0.29	0.03	-0.02	0.16	-0.19	1.14	0.78	0.78	1.07	0.96
Error g_a	0.00	0.01	0.61	3.69	2.72	1.32	2.66	0.22	0.04	1.17
Error g_b	0.09	0.00	0.00	0.03	0.04	0.02	0.05	0.05	0.00	0.00
									12.44	0.28

- ▶ Clearly $g_b(\mathbf{x})$ has a lower SSE.

Linear Classifiers via Multiple Regression

- ▶ There is always a single line that minimizes the SSE
- ▶ With a little matrix algebra, it is simple to derive the formula for this line.
- ▶ To fit model

$$\hat{y} = w_0 + w_1x_1 + \cdots + w_mx_m$$

- ▶ We can re-write this in matrix form as $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$, where

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_m \end{bmatrix} \quad \mathbf{X} = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1m} \\ 1 & x_{21} & x_{22} & \cdots & x_{2m} \\ \vdots & \vdots & \vdots & & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{nm} \end{bmatrix}$$

Least Squares Linear Classifier

- ▶ The model errors are defined by $\mathbf{e} = \mathbf{y} - \hat{\mathbf{y}}$
- ▶ The sum of squared errors is given by

$$SSE = \sum (y_i - \hat{y}_i)^2 = (\mathbf{y} - \mathbf{X}\mathbf{w})^2 = (\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

- ▶ Differentiate and set to zero

$$\frac{\partial(SSE)}{\partial \mathbf{w}} = -2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = 0$$

$$\mathbf{X}^T\mathbf{y} - \mathbf{X}^T\mathbf{X}\mathbf{w} = 0$$

$$\mathbf{X}^T\mathbf{X}\mathbf{w} = \mathbf{X}^T\mathbf{y}$$

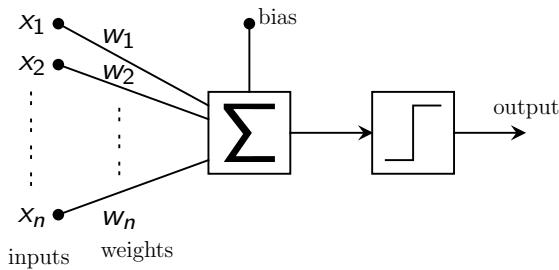
$$\mathbf{w} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$$

Assuming the inverse exists.

- ▶ Equivalent to Fisher's linear discriminant!

Linear Classifiers: Perceptron Approach

One of the first attempts to make an artificial neural network².



A *Perceptron* computes the sum of its weighted inputs and passes the result to a hard-limit threshold function.

Given input attributes x_1, x_2, \dots, x_m and response variable t that can take values -1 or $+1$, let

$$y = \psi\{w_0 + w_1x_1 + \dots + w_mx_m\} \quad \text{where} \quad \psi\{z\} = \begin{cases} +1 & \text{if } z \geq 0 \\ -1 & \text{if } z < 0 \end{cases}$$

²Rosenblatt, Frank (1958), The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain, Cornell Aeronautical Laboratory, Psychological Review, v65, No. 6, pp. 386-408.

Perceptron Training Rule

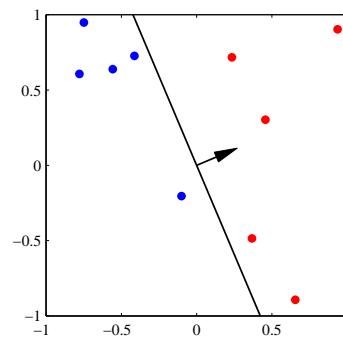
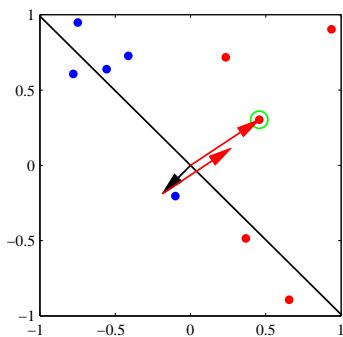
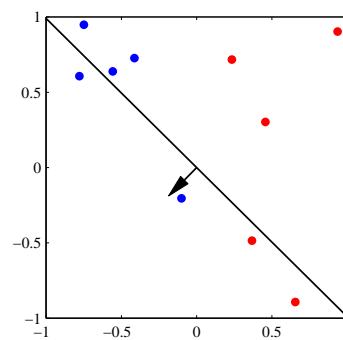
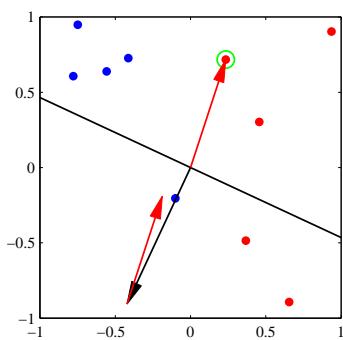
Algorithm PerceptronTraining(DataSet X , t)
Returns LinearModel w

```
initialise  $w$  to random values
initialise learning rate  $\eta$ 
do
  for  $i=1$  to  $n$ 
     $y_i = \psi(w, x_i)$ 
    for  $j=1$  to  $m$ 
       $\Delta w_j \leftarrow 0.5\eta(t_i - y_i)x_{ij}$ 
       $w_j \leftarrow w_j + \Delta w_j$ 
  while (Stopping( $t, y$ ) == false)
return  $w$ 
```

The Perceptron Rule

- ▶ Basic idea, cycle through training patterns, if a pattern is currently misclassified add/subtract the input vector to the weights.
 - ▶ This shifts the discriminant to make it more likely to classify that pattern correctly next time.
- ▶ η is the learning rate which moderates how much the weights are changed on each iteration.
 - ▶ The discriminant is unaffected by the magnitude of $\mathbf{w}^T \mathbf{x}$, so let $\eta = 1$.
- ▶ Stopping conditions vary, simplest is to stop when there is no change in \mathbf{y} or error is zero.
 - ▶ A complete pass through the dataset is made without modifying the weights.

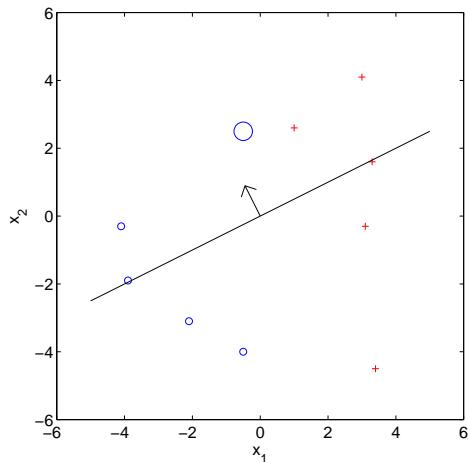
Example: The Perceptron Rule



Perceptron Example - Step #1

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} -1.0 \\ +2.0 \end{bmatrix}$$



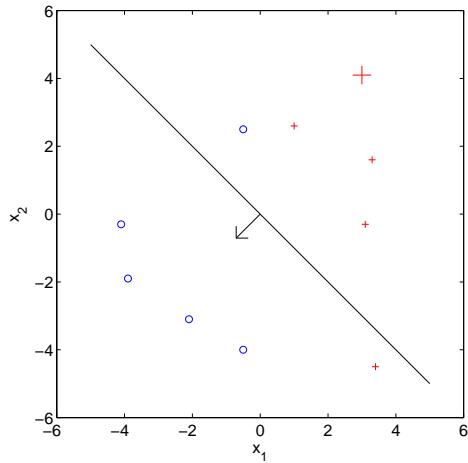
$$y_1 = \psi(-1.0 \times -0.5 + 2.0 \times 2.5) = \psi(5.5) \text{ -- WRONG!}$$

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i \\ &= [-1 \ 2] + 0.5 \times (-1 - +1)[-0.5 \ 2.5] = [-0.5 \ -0.5] \end{aligned}$$

Perceptron Example - Step #2

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} -0.5 \\ -0.5 \end{bmatrix}$$



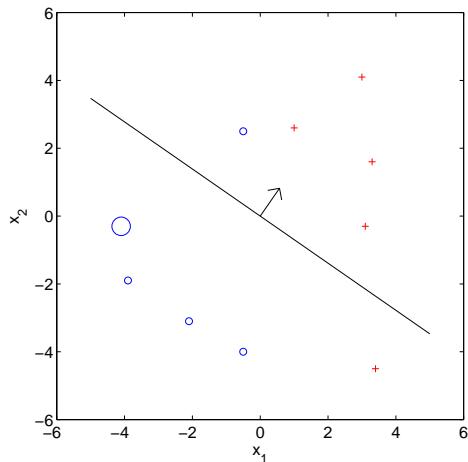
$$y_2 = \psi(-0.5 \times 3.0 + -0.5 \times 4.1) = \psi(-3.55) \text{ -- WRONG!}$$

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i \\ &= [-0.5 \ -0.5] + 0.5 \times (+1 - -1)[3.0 \ 4.1] = [2.5 \ 3.6] \end{aligned}$$

Perceptron Example - Step #3

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +2.5 \\ +3.6 \end{bmatrix}$$



$$y_3 = \psi(2.5 \times -4.1 + 3.6 \times -0.3) = \psi(-11.330000) \text{ -- RIGHT!}$$

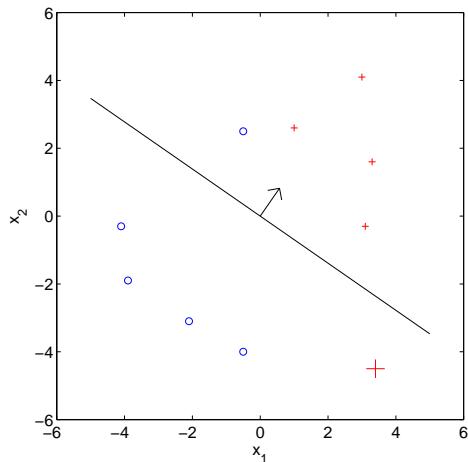
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [2.5 \ 3.6] + 0.5 \times (-1 - -1)[-4.1 \ -0.3] = [2.5 \ 3.6]$$

Perceptron Example - Step #4

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +2.5 \\ +3.6 \end{bmatrix}$$



$$y_4 = \psi(2.5 \times 3.4 + 3.6 \times -4.5) = \psi(-7.7) \text{ -- WRONG!}$$

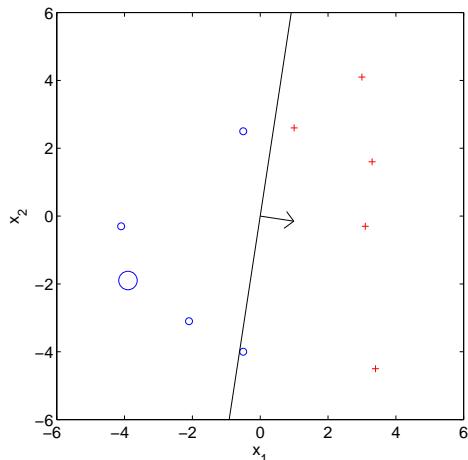
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [2.5 \ 3.6] + 0.5 \times (+1 - -1)[3.4 \ -4.5] = [5.9 \ -0.9]$$

Perceptron Example - Step #5

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +5.9 \\ -0.9 \end{bmatrix}$$



$$y_5 = \psi(5.9 \times -3.9 + -0.9 \times -1.9) = \psi(-21.3) \text{ -- RIGHT!}$$

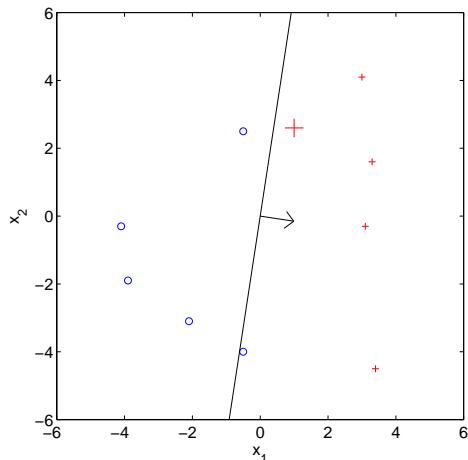
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [5.9 \quad -0.9] + 0.5 \times (-1 - -1)[-3.9 \quad -1.9] = [5.9 \quad -0.9]$$

Perceptron Example - Step #6

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +5.9 \\ -0.9 \end{bmatrix}$$



$$y_6 = \psi(5.9 \times 1.0 + -0.9 \times 2.6) = \psi(3.56) \text{ -- RIGHT!}$$

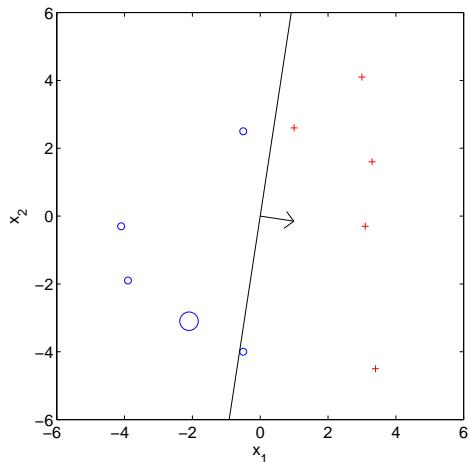
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [5.9 \quad -0.9] + 0.5 \times (+1 - +1)[1.0 \quad 2.6] = [5.9 \quad -0.9]$$

Perceptron Example - Step #7

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +5.9 \\ -0.9 \end{bmatrix}$$



$$y_7 = \psi(5.9 \times -2.1 + -0.9 \times -3.1) = \psi(-9.6) \text{ -- RIGHT!}$$

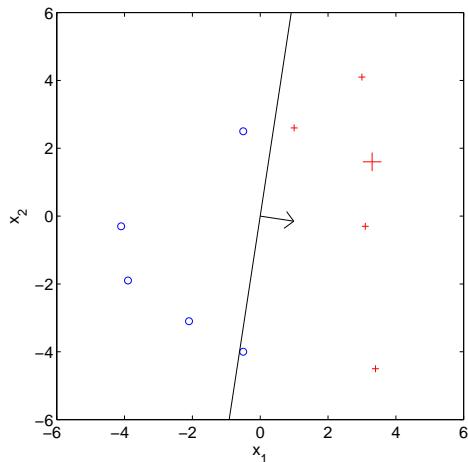
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [5.9 \quad -0.9] + 0.5 \times (-1 - -1)[-2.1 \quad -3.1] = [5.9 \quad -0.9]$$

Perceptron Example - Step #8

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +5.9 \\ -0.9 \end{bmatrix}$$



$$y_8 = \psi(5.9 \times 3.3 + -0.9 \times 1.6) = \psi(18.03) \text{ -- RIGHT!}$$

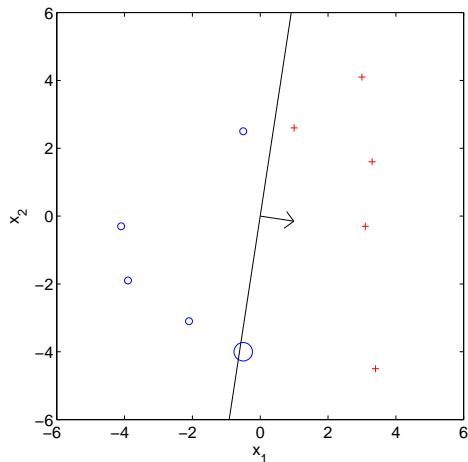
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [5.9 \quad -0.9] + 0.5 \times (+1 - +1)[3.3 \quad 1.6] = [5.9 \quad -0.9]$$

Perceptron Example - Step #9

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +5.9 \\ -0.9 \end{bmatrix}$$



$$y_9 = \psi(5.9 \times -0.5 + -0.9 \times -4.0) = \psi(0.65) \text{ -- WRONG!}$$

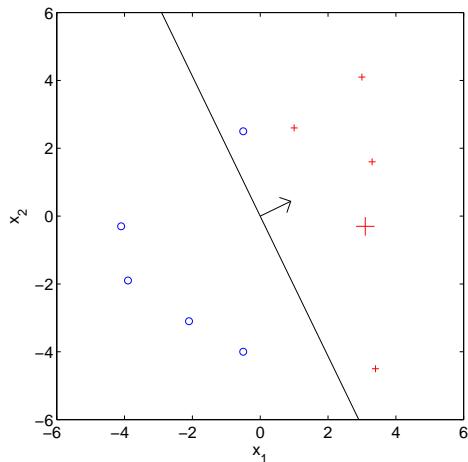
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [5.9 \ -0.9] + 0.5 \times (-1 - +1)[-0.5 \ -4.0] = [6.4 \ 3.1]$$

Perceptron Example - Step #10

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +6.4 \\ +3.1 \end{bmatrix}$$



$$y_{10} = \psi(6.4 \times 3.1 + 3.1 \times -0.3) = \psi(18.91) \text{ -- RIGHT!}$$

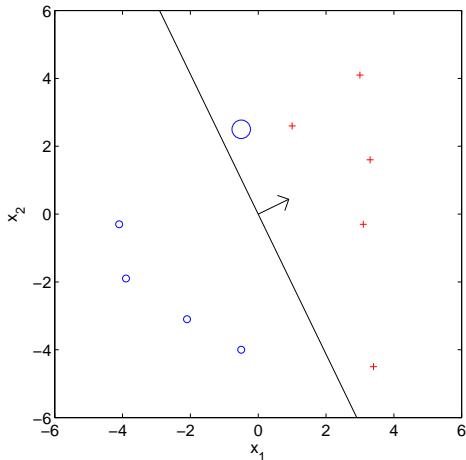
$$\mathbf{w} \leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i$$

$$= [6.4 \ 3.1] + 0.5 \times (+1 - +1)[3.1 \ -0.3] = [6.4 \ 3.1]$$

Perceptron Example - Step #11

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +6.4 \\ +3.1 \end{bmatrix}$$



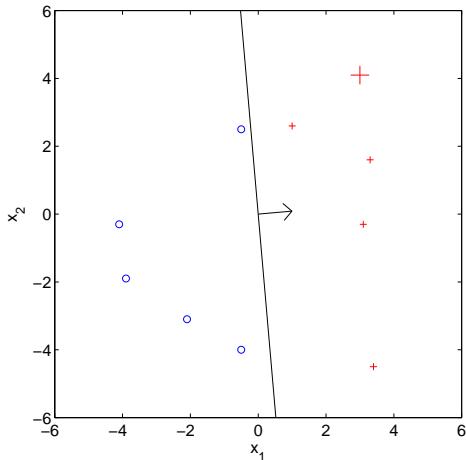
$$y_1 = \psi(6.4 \times -0.5 + 3.1 \times 2.5) = \psi(4.55) \text{ - WRONG!}$$

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i \\ &= [6.4 \ 3.1] + 0.5 \times (-1 - +1)[-0.5 \ 2.5] = [6.9 \ 0.6] \end{aligned}$$

Perceptron Example - Step #12

x_1	x_2	t
-0.5	+2.5	-1
+3.0	+4.1	+1
-4.1	-0.3	-1
+3.4	-4.5	+1
-3.9	-1.9	-1
+1.0	+2.6	+1
-2.1	-3.1	-1
+3.3	+1.6	+1
-0.5	-4.0	-1
+3.1	-0.3	+1

$$\mathbf{w} = \begin{bmatrix} +6.9 \\ +0.6 \end{bmatrix}$$



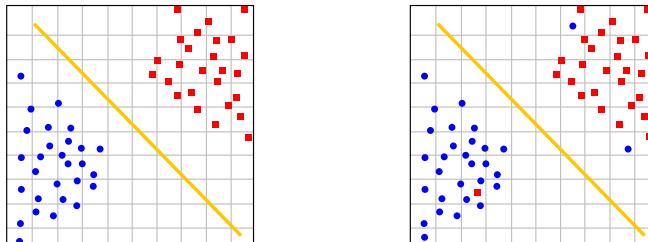
$$y_2 = \psi(6.9 \times 3.0 + 0.6 \times 4.1) = \psi(23.16) \text{ - RIGHT!}$$

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} + 0.5\eta(t_i - y_i)\mathbf{x}_i \\ &= [6.9 \ -0.6] + 0.5 \times (+1 - +1)[3.0 \ 4.1] = [6.9 \ 0.6] \end{aligned}$$

Linearly Separable Problems

- ▶ A problem is linearly separable if there exists a linear discriminant that gives perfect classification

Linearly Separable Not Linearly Separable



- ▶ The perceptron rule will converge to a correct solution for a linearly separable problem in a finite number of steps³
- ▶ It may **fail to converge** on linearly inseparable problems.

³Novikoff, A. B. (1962). On convergence proofs on perceptrons. Symposium on the Mathematical Theory of Automata, 12, 615-622. Polytechnic Institute of Brooklyn.

Off-Line Perceptron Algorithm

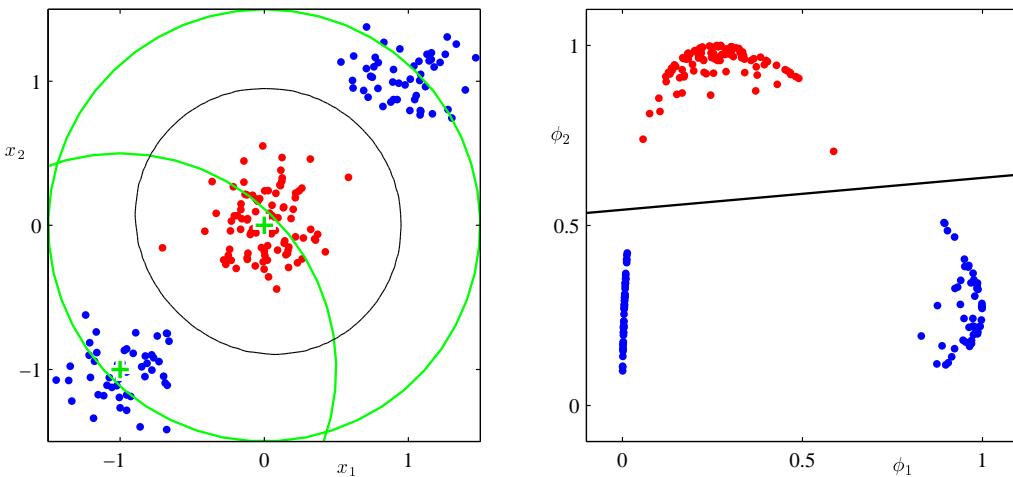
Algorithm GradientDescentTraining(DataSet X, t)
Returns Classifications R

```
initialise  $\mathbf{w}$  to random values
initialise learning rate  $\eta$ 
do
    initialise  $\Delta\mathbf{w}$  to zeros
    for i=1 to n
         $y_i = \psi(\mathbf{w}, \mathbf{x}_i)$ 
        for j=1 to m
             $\Delta w_j \leftarrow \Delta w_j + \frac{1}{2}\eta(t_i - y_i)x_{ij}$ 
        for j=1 to m
             $w_j \leftarrow w_j + \Delta w_j$ 
    while (Stopping( $\mathbf{r}, \mathbf{y}$ )==false)
    return  $\mathbf{w}$ 
```

Non-Linear Transformation

- ▶ Construct features by placing a Gaussian basis function on two patterns

$$\phi(\mathbf{x}) = \exp \{-\gamma \|\mathbf{x} - \mathbf{x}_a\|^2\}$$



Summary

- ▶ Linear models can be justified in many ways:
 - ▶ Optimal classifiers for Gaussian data.
 - ▶ Fisher's linear discriminant.
 - ▶ Least-squares regression.
 - ▶ The Perceptron.
- ▶ Advantages:
 - ▶ Provably optimal for data from [some] normal distributions
 - ▶ Practically effective for a wide range of problems
 - ▶ Not particularly sensitive to redundant features
- ▶ Disadvantages:
 - ▶ Cannot solve even simple non-linear problems
 - ▶ Sensitive to correlated features
 - ▶ Sensitive to ill conditioned problems (matrix inverse)